

ATTACKING RUST

MODERN SYSTEM APPLICATION
VULNERABILITIES COMPREHENSIVE
ANALYSIS



Attacking Rust

• May 13, 2024 •  14 min read

Table of contents

Cargo Dependency Confusing

Unsafe Code Usage

Integer Overflow

Panics in Rust Code

memory leaks

Uninitialized memory

Foreign Function Interface

OOB Read plus

race condition to escalate privileges

TOCTAU race condition

out-of-bounds array access

References

Show less 

"Attacking Rust" delves into the intricacies of identifying and mitigating security vulnerabilities within Rust codebases. Despite Rust's reputation for strong memory safety and thread concurrency, no programming language is immune to potential exploits. This article navigates through common attack vectors such as buffer overflows, race conditions, and injection attacks, illustrating how they can manifest

within Rust applications. By examining these examples through the lens of real-world examples and discussing their implications, developers gain valuable insights into fortifying their Rust projects against potential threats.

Moreover, "Attacking Rust" goes beyond mere identification of vulnerabilities; it equips developers with the tools and techniques necessary for proactive defense. From leveraging Rust's ownership system to prevent data races to employing secure coding patterns and libraries, this article empowers developers to construct robust and resilient software. By adopting a proactive stance towards security, developers can not only safeguard their applications from exploitation but also contribute to fortifying the broader Rust ecosystem against emerging threats.

Cargo Dependency Confusing

In the Rust ecosystem, Cargo serves as the indispensable tool for managing dependencies, building projects, and executing various development tasks. With commands like `cargo build`, `cargo test`, and `cargo doc`, developers can seamlessly compile their code, run tests, and generate documentation. However, the convenience offered by Cargo also introduces potential security risks, particularly concerning dependency management and configuration customization.

One area of concern is the reliance on external dependencies fetched through Cargo. While Cargo downloads packages over HTTPS, it does not validate the registry index, leaving room for malicious actors to inject compromised dependencies. Although efforts are made to ensure the security of crates.io and its associated GitHub repository, vulnerabilities may still arise. Developers should remain vigilant and consider alternative methods for dependency installation in critical scenarios.

Moreover, Cargo's flexibility in customizing build configurations through the `Cargo.toml` file can inadvertently introduce vulnerabilities. For instance, overriding default options related to debug assertions and integer overflow checks in development profiles can lead to undetected bugs, compromising the application's security. Adherence to established rules, such as refraining from overriding certain

options in development to ensure consistency, integrity and robustness.

COPY 

```
[dependencies]
rand = "0.8.4"
```

The code specifies a dependency on the "rand" crate without pinning it to a specific version. This lack of version specification can lead to dependency confusion, where Cargo might fetch a different version of the "rand" crate than intended, potentially introducing security vulnerabilities or unexpected behavior into the project.

We specify the exact version of the "rand" crate ("0.8.4") that our project depends on. This pinning ensures that Cargo fetches the intended version of the dependency, reducing the risk of dependency confusion and helping to maintain the security and stability of the project.

In addition, Cargo provides mechanisms for altering its behavior using environment variables like `RUSTC`, `RUSTC_WRAPPER`, and `RUSTFLAGS`. While these options offer flexibility, their indiscriminate use can undermine the predictability and reliability of the build process. Centralizing compiler options and flags within the `Cargo.toml` configuration file is recommended to ensure consistency and mitigate the risk of unintended consequences. By adhering to best practices and exercising caution when configuring Cargo, developers can bolster the security posture of their Rust projects and safeguard against potential attacks.

Unsafe Code Usage

Rust's safety guarantees rely heavily on its type system and ownership model to prevent common pitfalls like memory overflows, null pointer dereferences, and data races. However, the language also provides an escape hatch in the form of the `unsafe` keyword, allowing developers to bypass these safety checks when necessary.

While `unsafe` blocks are sometimes unavoidable when interfacing with low-level code or accessing hardware resources, they also introduce risks that can compromise the integrity and security of Rust programs.

Attack Scenario: Consider a scenario where a developer utilizes `unsafe` blocks without proper justification or understanding of the associated risks. For example, they might use `unsafe` to perform memory manipulations or system calls without implementing adequate safeguards, potentially exposing the application to memory corruption vulnerabilities or undefined behavior. Additionally, improper handling of `unsafe` code in libraries or APIs could propagate vulnerabilities to downstream consumers, leading to widespread security issues.

Rust's safety features are a key aspect of its design, aimed at preventing memory safety vulnerabilities such as buffer overflows and null pointer dereferences. However, Rust provides the `unsafe` keyword as a way to bypass these safety checks when necessary, typically for interfacing with low-level code or accessing system resources. While `unsafe` blocks are sometimes unavoidable, their misuse can lead to memory corruption, undefined behavior, and other security vulnerabilities.

```
unsafe fn unchecked_access(ptr: *mut i32) {
    *ptr = 42;
}

fn main() {
    let mut data: i32 = 0;
    unsafe {
        unchecked_access(&mut data);
    }
    println!("Data: {}", data);
}
```

COPY 

In this non-compliant example, the `unsafe` block is marked as `unsafe`, allowing direct manipulation of memory without Rust's safety checks. This code is potentially dangerous as it can lead to memory corruption if `ptr` is not properly validated or managed.

Compliant Rust Code:

```
fn safe_access(ptr: &mut i32) {
    *ptr = 42;
}

fn main() {
    let mut data: i32 = 0;
    safe_access(&mut data);
    println!("Data: {}", data);
}
```

COPY 

In the compliant example, the `safe_access` function operates on a mutable reference (`&mut i32`), ensuring that Rust's safety guarantees are upheld. By avoiding the use of `unsafe` blocks and leveraging safe Rust constructs, the code remains secure and free from potential memory safety vulnerabilities.

Integer Overflow

Integer overflows occur when the result of an arithmetic operation exceeds the maximum value that can be represented by the data type. In Rust, integer overflows can lead to unexpected behavior, security vulnerabilities, and even application crashes. While Rust provides some safeguards against integer overflows, developers must remain vigilant and adopt appropriate mitigation strategies to prevent these vulnerabilities.

```
fn vulnerable_operation(x: u8) -> u8 {
    x + 200
}

fn main() {
    let result = vulnerable_operation(50);
    println!("Result: {}", result);
}
```

In this non-compliant example, the `vulnerable_operation` function performs an arithmetic operation (`x + 200`) without considering the possibility of integer overflow. This code may produce unexpected results or even panic in debug mode, leading to potential security vulnerabilities.

```
use std::num::Wrapping;

fn safe_operation(x: u8) -> u8 {
    let result = Wrapping(x) + Wrapping(200);
    result.0
}

fn main() {
    let result = safe_operation(50);
    println!("Result: {}", result);
}
```

COPY 

In the compliant example, the `safe_operation` function uses the `Wrapping` type to perform the arithmetic operation in a safe and predictable manner. By wrapping the values with `Wrapping` before performing the addition, the code explicitly handles potential integer overflows, ensuring that the result is well-defined and free from vulnerabilities.

Panics in Rust Core



Panics occur when Rust encounters unrecoverable errors or violations of invariants during program execution. While panics provide a mechanism for handling exceptional conditions, excessive or uncontrolled use of panics can lead to unpredictable behavior, application crashes, and security vulnerabilities. It is essential for Rust developers to employ proper error handling mechanisms, such as `Result` or `Option`, instead of relying on panics to handle errors.

COPY

```
fn divide(a: i32, b: i32) -> i32 {
    if b == 0 {
        panic!("Division by zero");
    }
    a / b
}

fn main() {
    let result = divide(10, 0);
    println!("Result: {}", result);
}
```

In this non-compliant example, the `divide` function panics when attempting to divide by zero. Panics should be avoided in favor of returning a `Result` to indicate error conditions.

COPY

```
fn divide(a: i32, b: i32) -> Result<i32, &'static str> {
    if b == 0 {
        Err("Division by zero")
    } else {
        Ok(a / b)
    }
}
```



```
fn main() {
    match divide(10, 0) {
        Ok(result) => println!("Result: {}", result),
        Err(err) => println!("Error: {}", err),
    }
}
```

In the compliant example, the `divide` function returns a `Result` type instead of panicking. This allows the caller to handle the error condition gracefully and take appropriate action. By adopting proper error handling practices, the code becomes more robust and resistant to unexpected failures.


memory leaks

Memory leaks occur when a program fails to release memory that is no longer needed, leading to excessive memory consumption and potential performance degradation. In Rust, memory leaks can occur when resources are not properly managed, particularly when using manual memory management functions such as `forget`, `Box::leak`, or `Box::into_raw`. These functions bypass Rust's memory safety guarantees and can result in unreclaimed memory and resource leaks.

```
use std::mem;

fn create_resource() -> String {
    String::from("Resource")
}

fn main() {
    let resource = create_resource();
    mem::forget(resource);
}
```

COPY 

In this non-compliant example, the `create_resource` function creates a `String` resource, which is then assigned to a variable. This results in a memory leak as the allocated memory for the `String` is never reclaimed, leading to potential resource exhaustion and performance issues.

```
fn create_resource() -> String {
    String::from("Resource")
}

fn main() {
    let _resource = create_resource();
}
```

COPY 

In the compliant example, the `create_resource` function returns a `String` resource, which is then used within the `main` function. The resource is automatically deallocated when it goes out of scope, ensuring that memory is reclaimed properly and preventing memory leaks.

Uninitialized memory

Uninitialized memory occurs when a program attempts to read or write data from memory locations that have not been properly initialized with valid values. In Rust, uninitialized memory can be inadvertently introduced through the use of deprecated functions like `std::mem::uninitialized` or the `MaybeUninit` type. Accessing uninitialized memory can lead to memory safety issues, data corruption, and potential security vulnerabilities.

```
use std::mem;

fn process_uninitialized() -> u32 {
    let mut value: u32;
```

COPY 

```

unsafe {
    value = mem::uninitialized();
}

// Perform operations using uninitialized value
value

}

fn main() {
    let result = process_uninitialized();
    println!("Result: {}", result);
}

```

In this non-compliant example, the `process_uninitialized` function uses `std::mem::uninitialized` to create an uninitialized `u32` value, which is then used without initialization. This code can lead to undefined behavior and potential security vulnerabilities.

```

fn process_initialized() -> u32 {
    let value: u32 = 42; // Initialize value to a default or known
    value
    // Perform operations using initialized value
    value
}

fn main() {
    let result = process_initialized();
    println!("Result: {}", result);
}

```

In the compliant example, the `process_initialized` function initializes the `u32` value to a default or known value before performing operations. By ensuring that all variables are properly initialized before use, the code remains memory safe and free from uninitialized memory vulnerabilities.

Foreign Function Interface



The Foreign Function Interface (FFI) in Rust allows seamless interoperability between Rust and other languages, particularly C. While FFI enables powerful capabilities, it also introduces security risks, as it bypasses Rust's safety guarantees and exposes the program to potential vulnerabilities such as memory safety issues, data corruption, and undefined behavior.

COPY

```
// Export a C-compatible function
#[no_mangle]
unsafe extern "C" fn mylib_f(param: u32) -> i32 {
    if param == 0xCAFEBADE {
        0
    } else {
        -1
    }
}

fn main() {
    let x = -1;
    let result = unsafe { mylib_f(x) };
    println!("Result: {}", result);
}
```

In this non-compliant example, the Rust code exposes a C-compatible function `mylib_f` to be called from C code. However, the function does not perform proper parameter validation, potentially leading to security vulnerabilities if called with invalid arguments.

COPY

```
// Export a C-compatible function with proper parameter validation
#[no_mangle]
unsafe extern "C" fn mylib_f(param: u32) -> i32 {
```

```

    if param == 0xCAFEBADE {
        0
    } else {
        -1
    }
}

fn main() {
    let x = 0xCAFEBADE; // Valid parameter value
    let result = unsafe { mylib_f(x) };
    println!("Result: {}", result);
}

```

In the compliant example, the Rust code properly validates the parameter passed to the C-compatible function `mylib_f` before performing any operations. By ensuring that all inputs are validated and sanitized, the code mitigates the risk of potential security vulnerabilities through FFI.

OOB Read plus

Out-of-bounds (OOB) read vulnerabilities occur when a program attempts to read data from memory locations outside the bounds of an allocated buffer. In Rust, OOB read vulnerabilities can lead to information disclosure, memory corruption, and potential security exploits. These vulnerabilities typically arise from improper bounds checking or unchecked array accesses.

```

fn oob_read_vulnerability(data: &[u8], index: usize) -> u8 {
    data[index]
}

fn main() {
    let data = vec![1, 2, 3, 4, 5];
    let result = oob_read_vulnerability(&data, 10);
}

```

COPY 

```
println!("Result: {}", result);
```

```
}
```



In this non-compliant example, the `oob_read_vulnerability` function reads data from a slice without proper bounds checking. When called with an out-of-bounds index, the function accesses memory beyond the bounds of the slice, leading to potential OOB read vulnerabilities.

COPY

```
fn safe_read(data: &[u8], index: usize) -> Option<u8> {
    if index < data.len() {
        Some(data[index])
    } else {
        None
    }
}

fn main() {
    let data = vec![1, 2, 3, 4, 5];
    if let Some(result) = safe_read(&data, 10) {
        println!("Result: {}", result);
    } else {
        println!("Index out of bounds");
    }
}
```

In the compliant example, the `safe_read` function performs bounds checking to ensure that the index is within the bounds of the data slice before attempting to access the element. By returning an `Option<u8>` instead of directly accessing the data, the code prevents OOB read vulnerabilities and gracefully handles out-of-bounds access scenarios.

race condition to escalate privileges

A race condition occurs when the outcome of a program depends on the timing or sequence of events that are not under the program's control. In Rust, race conditions can lead to security vulnerabilities, especially when they are exploited to escalate privileges or gain unauthorized access to sensitive resources. Privilege escalation via race conditions typically involves manipulating shared resources in a concurrent environment to bypass access controls or gain elevated privileges.

COPY 

```
use std::sync::{Arc, Mutex};
use std::thread;

fn insecure_increment(counter: Arc<Mutex<u32>>) {
    let mut value = counter.lock().unwrap();
    *value += 1;
}

fn main() {
    let counter = Arc::new(Mutex::new(0));
    let mut threads = vec![];

    for _ in 0..10 {
        let counter_clone = Arc::clone(&counter);
        let handle = thread::spawn(move || {
            insecure_increment(counter_clone);
        });
        threads.push(handle);
    }

    for handle in threads {
        handle.join().unwrap();
    }

    println!("Final Counter Value: {:?}", *counter.lock().unwrap());
}
```

In this non-compliant example, multiple threads can access and modify a shared counter value without synchronization, which can result in a race condition where the final counter value is non-deterministic and may not reflect the actual number of increments performed. An attacker could potentially exploit this race condition to manipulate the counter value and escalate privileges.

```
use std::sync::{Arc, Mutex};
use std::thread;

fn secure_increment(counter: Arc<Mutex<u32>>) {
    let mut value = counter.lock().unwrap();
    *value += 1;
}

fn main() {
    let counter = Arc::new(Mutex::new(0));
    let mut threads = vec![];

    for _ in 0..10 {
        let counter_clone = Arc::clone(&counter);
        let handle = thread::spawn(move || {
            secure_increment(counter_clone);
        });
        threads.push(handle);
    }

    for handle in threads {
        handle.join().unwrap();
    }

    println!("Final Counter Value: {:?}", *counter.lock().unwrap());
}
```

COPY 

In the compliant example, `lock` ensures that only one thread can access the `counter`, preventing race conditions and ensuring the integrity of the counter. By employing thread-safe synchronization mechanisms, the code mitigates the risk of privilege escalation through race conditions.

TOCTAU race condition

Time-of-Check to Time-of-Use (TOCTOU) race conditions occur when a program's behavior depends on the state of a resource at two separate points in time: the time of checking (TOC) and the time of use (TOU). In Rust, TOCTOU vulnerabilities can lead to security exploits, especially when they are exploited to manipulate resource states between the check and the use, potentially bypassing security checks or violating access controls.

COPY 

```
use std::fs;

fn insecure_file_access(filename: &str) -> Result<(), std::io::Error>
{
    // Check if the file exists
    if fs::metadata(filename).is_ok() {
        // Attempt to read the file
        let _contents = fs::read_to_string(filename)?;
        Ok(())
    } else {
        Err(std::io::Error::new(
            std::io::ErrorKind::NotFound,
            "File not found",
        ))
    }
}

fn main() {
    let filename = "example.txt";
```

```

    if let Err(err) = insecure_file_access(filename) {
        eprintln!("{}", err);
    }
}

```

In this non-compliant example, the `insecure_file_access` function first checks if a file exists using `fs::metadata`, then attempts to read the file using `fs::read_to_string`. However, there is a window of opportunity between the check and the use where the file's state may change, leading to a TOCTOU vulnerability. An attacker could potentially manipulate the file's state between the check and the use to bypass access controls or read sensitive data.

```

use std::fs;

fn secure_file_access(filename: &str) -> Result<(), std::io::Error> {
    // Attempt to read the file
    let _contents = fs::read_to_string(filename)?;
    Ok(())
}

fn main() {
    let filename = "example.txt";
    if let Err(err) = secure_file_access(filename) {
        eprintln!("Error: {}", err);
    }
}

```

In the compliant example, the `secure_file_access` function directly attempts to read the file using `fs::read_to_string` without prior checking. By avoiding the separate check and use steps, the code eliminates the window of opportunity for TOCTOU vulnerabilities. Instead, any errors encountered during file access are directly handled, ensuring the integrity and security of the operation.

out-of-bounds array access



Out-of-bounds array access occurs when a program attempts to read or write to memory locations outside the bounds of an allocated array. In Rust, out-of-bounds array access can lead to memory corruption, data leakage, and potentially exploitable security vulnerabilities. These vulnerabilities typically arise from improper bounds checking or unchecked array accesses.

COPY

```
fn non_compliant_out_of_bounds_access() {  
    let array = [1, 2, 3, 4, 5];  
    let index = 10; // Accessing index out of bounds  
    let value = array[index];  
    println!("Value at index {}: {}", index, value);  
}
```

In this non-compliant example, the `non_compliant_out_of_bounds_access` function attempts to access an element of the `array` at an index (`index`) that is out of bounds. This can lead to undefined behavior, memory corruption, or potential security vulnerabilities if sensitive data is accessed or overwritten.

COPY

```
fn compliant_out_of_bounds_access() {  
    let array = [1, 2, 3, 4, 5];  
    let index = 2; // Accessing within bounds  
    if let Some(value) = array.get(index) {  
        println!("Value at index {}: {}", index, value);  
    } else {  
        println!("Index {} is out of bounds", index);  
    }  
}
```

In the compliant example, the `get` function uses the `get` method to access elements. This ensures that the program does not attempt to access elements outside the bounds of the array, preventing out-of-bounds array access vulnerabilities. Additionally, the code gracefully handles the case where the index is out of bounds by checking the result of `array.get(index)` and printing an error message.

References

- <https://redteamguides.com/>
- <https://anssi-fr.github.io/rust-guide/>

Subscribe to our newsletter

Read articles from **DevSecOpsGuides** directly inside your inbox. Subscribe to the newsletter, and don't miss out.

reza.rashidi.business@gmail.com

SUBSCRIBE

Rust

Devops

Developer


DevSecOps

appsec

Written by



Reza Rashidi

 Add your bio

Published on



DevSecOpsGuides

Add blog description

MORE ARTICLES

RR

Reza Rashidi



Attacking Java

Attacking Java applications requires a nuanced understanding of both the language's intricacies and ...

RR

Reza Rashidi



Attacking PHP

In modern PHP applications, attackers exploit various vulnerabilities to compromise systems, steal d...

RR

Reza Rashidi



Attacking NodeJS Application

When it comes to securing Node.js applications, understanding potential attack vectors is paramount....

©2024 DevSecOpsGuides

[Archive](#) · [Privacy_policy](#) · [Terms](#)



Write on Hashnode

Powered by [Hashnode](#) - Home for tech writers and readers